

PRG One Workspace progress and next steps

Prepared by Varchasv Gupta for project stakeholders
Status as at 3 October 2026
Public project summary with confidential operational detail excluded

The current position

The project has progressed from an interview demonstration to a working, privately hosted prototype. It brings a case's next action, enrichment evidence, review decisions, notes and handoff into one place. The intended benefit is less switching between screens and less repeated administration while keeping JobAdder as the main recruitment record.

The workspace functions are built and tested. Live platform connections and a real operational pilot are still to be completed. It can save fictional cases and reviews today. It does not yet fetch live platform data, update JobAdder or send messages. No actual time saving, cash saving or improvement in recruitment outcomes has been demonstrated.

The next decision is to agree one small enrichment pilot, name its business and system owners, and arrange approved access. A focused pilot will show whether to configure existing tools or build a small additional connection.

What stakeholders can see now

Capability	Current position	Practical purpose
Dashboard and work board	Built with saved cases, due dates and priorities	Make the next action visible.
Three workflow types	Permanent, contract and client research prototypes	Keep different tasks and handoffs separate.
Enrichment review	Current and proposed values, evidence, accept, hold and reject	Make changes inspectable before use.
Case notes and history	Saved notes, stage checkpoints and change history	Retain context through a handoff.
Handoff exports	Review packets and text drafts	Prepare information for manual review.
Platform hub and search	23 platform entries, shortcuts and setup notes	Find the right tool and understand its role.
Cost and pilot records	Saved assumptions and aggregate observations	Test value without presenting assumptions as results.
Live integrations and live AI	Not configured	Require approved access and further implementation.

What a fully working process should look like

End to end means a permitted record enters the workflow, useful changes are proposed, a person reviews them, approved changes reach the correct system, and the outcome can be measured and supported.

1. **Choose the right records.** A small, agreed group is selected from the recruitment system. Each record has a stable ID so the same person or contact is not confused with another.
2. **Find a useful missing or outdated value.** Check agreed fields only. An approved source supplies the evidence and its date. Existing tools should be used where they already provide this capability.
3. **Prepare a suggested change.** Show the current value, proposed value and supporting evidence. A suggestion must not silently become an accepted fact.
4. **Review the change.** A named reviewer accepts, holds or rejects it. Missing evidence and conflicting information remain visible.
5. **Apply the approved update.** Before changing JobAdder, check whether the original value has changed since review. Record what was updated and avoid duplicate writes.
6. **Carry the context forward.** Make the next action, receiving owner and relevant source references available. Any outward message or automation needs its own approved trigger and permissions.
7. **Measure and support the result.** Record useful updates, errors, review effort and total cost. Someone must own failures, corrections and ongoing support.

How far along that process is the project

Evidence of progress: 45 automated tests passed for the current private prototype. Desktop and mobile screens were checked. Tests cover permissions, competing edits, stage checks and evidence handling. This is evidence of prototype behaviour, not production readiness or a live integration audit.

The current prototype supports manual preparation, field review, saved coordination and handoff preparation using fictional data. The live intake, provider connection, system update and operational monitoring stages remain to be implemented and tested.

Platform links are shortcuts, not data connections. A readiness checkbox records planning progress; it does not prove that access works. Assignee names are labels, not invitations or shared access. A completed checklist records a person's assertions; it does not independently verify compliance or hiring suitability.

The helper that finds skill phrases uses fixed matching rules. There is no live generative AI assistant or enrichment provider running in this version. Adding one is a separate decision based on need, evidence quality, permitted data use and cost.

Scope recommendation: prove one useful enrichment journey first. Avoid asking consultants to maintain a second version of their whole recruitment pipeline in this workspace.

Information and decisions needed from PRG

The following is a request list for the project team. Suggested owners are roles to nominate, not people already assigned. Supply sensitive information through PRG-approved private channels; do not put it in GitHub issues, repository files or ordinary chat messages.

Requested item	What a useful answer contains	Suggested owner
Pilot purpose and boundary	One desk or cohort, record type, fields to improve, exclusions and intended business result.	Sponsor and desk lead
Record and field definitions	System of record, stable IDs, exact field names, allowed values, custom fields and duplicate rules.	ATS owner
Review rules	Approved sources, freshness rules, who may approve, conflict handling and fields that must never be inferred.	Process owner and reviewer
Existing capabilities	Current modules and licences, enabled native integrations, limits, field coverage and configuration owner.	Tool owners
Safe access for testing	Approved sandbox or test tenant, test accounts, least-privilege access and a secure credential setup route.	IT and platform owners
Example inputs and outputs	Fictional or approved redacted examples of a current record, source evidence, expected update and failure case.	Desk lead and ATS owner
Data handling approval	Approved hosting and providers, permitted data, retention and deletion rules, access roles and incident contact.	IT and privacy owner
Handoff requirements	Receiving role, necessary fields, approved event trigger, permitted recipients and exception handling.	Operations and process owner
Costs and baseline	Relevant invoices, contract limits and renewals, usage, task frequency and observed manual effort.	Finance and desk lead
Acceptance and ownership	Agreed quality and cost thresholds, rollout approver, support owner and escalation process.	Sponsor and IT

The minimum starting pack

For the first connection, the project needs a written pilot boundary, an ATS contact, an approved testing route, a field map, fictional or approved test examples, and named reviewers. An API is a controlled way for systems to exchange data; permission to use a website does not necessarily include API access.

Actual passwords, API keys, candidate records, private system links, contracts and invoices should stay in approved private systems. This public document lists what is needed without disclosing those details.

Recommended first access: a narrowly scoped, read-only connection. This lets the team verify record matching and field interpretation before enabling updates.

Remaining delivery stages

The sequence below is proposed. Dates and effort estimates should be agreed after access, scope and vendor constraints are known. Progress should be reported against working evidence rather than an unsupported percentage complete.

Stage	Work to complete	Evidence needed before moving on
1 Agree the pilot	Confirm the starting pack and compare existing-tool capability with the gap.	Written scope, owners, data approval and test access.
2 Connect for reading	Map IDs and fields; show where information came from and when it was refreshed.	Correct matching on agreed test records; visible access and sync failures.
3 Prepare and review	Connect one approved source or existing capability to suggested changes.	Traceable suggestions; held conflicts; measured reviewer effort.
4 Test controlled updates	Apply approved changes in a sandbox with current-value checks, duplicate protection and correction procedures.	Successful updates, blocked conflicting updates, audit trail and recovery tests.
5 Run a small live pilot	Use an approved cohort with supervision and a clear stop condition.	Agreed quality, usefulness, effort and cost results, including failed attempts.
6 Hand over and expand	Document support, permissions, monitoring, training and any further integrations.	Named operating owner, runbook, recovery exercise and sponsor sign-off.

Improvements that would make adoption easier

Prioritise a short guided first run, clearer save indicators, warnings before losing unsaved input, a readable handoff report and automatic linking when a case starts from an enrichment record. Add saved views for pending reviews and overdue work, followed by controlled case archiving and reopening.

Shared assignments, notifications and live sync status need further implementation. They should reflect the agreed operating model rather than appear connected before they are ready.

What fully functional must mean before rollout

A normal case must complete from intake to a verified update. An exception case must stop safely and show the owner what to do. The team must also demonstrate correct permissions, current-value checks, no duplicate updates, source traceability, recovery from a failed connection and clear support ownership.

The sponsor should approve measurable quality and cost thresholds before the pilot. No arbitrary accuracy target or delivery date is committed by this document.

How progress and value will be reported

Each update should state what changed, what evidence supports it, what remains blocked, which decision is needed and the next checkable milestone. Separate a feature being built from it being tested in a sandbox, trialled with approved live data and accepted for ongoing use.

Measure	Plain English meaning	Why it matters
Useful coverage	Records receiving a useful verified update divided by records attempted.	Avoids counting empty or irrelevant results as success.
Audited accuracy	Correct suggested fields divided by suggested fields independently checked.	Shows quality. Zero audited fields means unknown accuracy.
Review effort	Total checking, correction and rework time.	Prevents reporting gross automation time as the net benefit.
Cost per useful record	Total pilot cost divided by records with useful verified updates.	Includes failed attempts and support effort.
Update reliability	Approved updates applied correctly, with failed or duplicate attempts tracked.	Becomes meaningful once live updates are implemented.
Net business effect	Avoidable cash cost and time released, reported separately.	Staff capacity is not automatically a payroll saving.

Cost reduction decisions

Reuse included capabilities and process only useful records and fields. Compare native integrations before adding subscriptions or custom connections. Check usage, contracts and downstream dependencies before cancelling a tool.

Actual savings are not established. Measure manual and assisted effort, provider usage, support, transition costs and genuinely removable spend. Report negative results as well as positive ones.

Decisions requested at the next stakeholder meeting

1. Select one pilot use case and the fields it will cover.
2. Nominate the sponsor, ATS owner, reviewer and support contact.
3. Agree the approved test environment and data handling route.
4. Confirm whether existing capabilities can cover the proposed pilot.
5. Agree acceptance measures and the evidence required for a live rollout decision.

Instructions: Read the [user guide](#), then complete the [implementation checklist](#) in a private working copy. This editable Markdown brief also has a [printable PDF](#).

Repository boundary: Public code remains the earlier portfolio version. The current workbench source and operational configuration are private. Site access is separate from repository access.